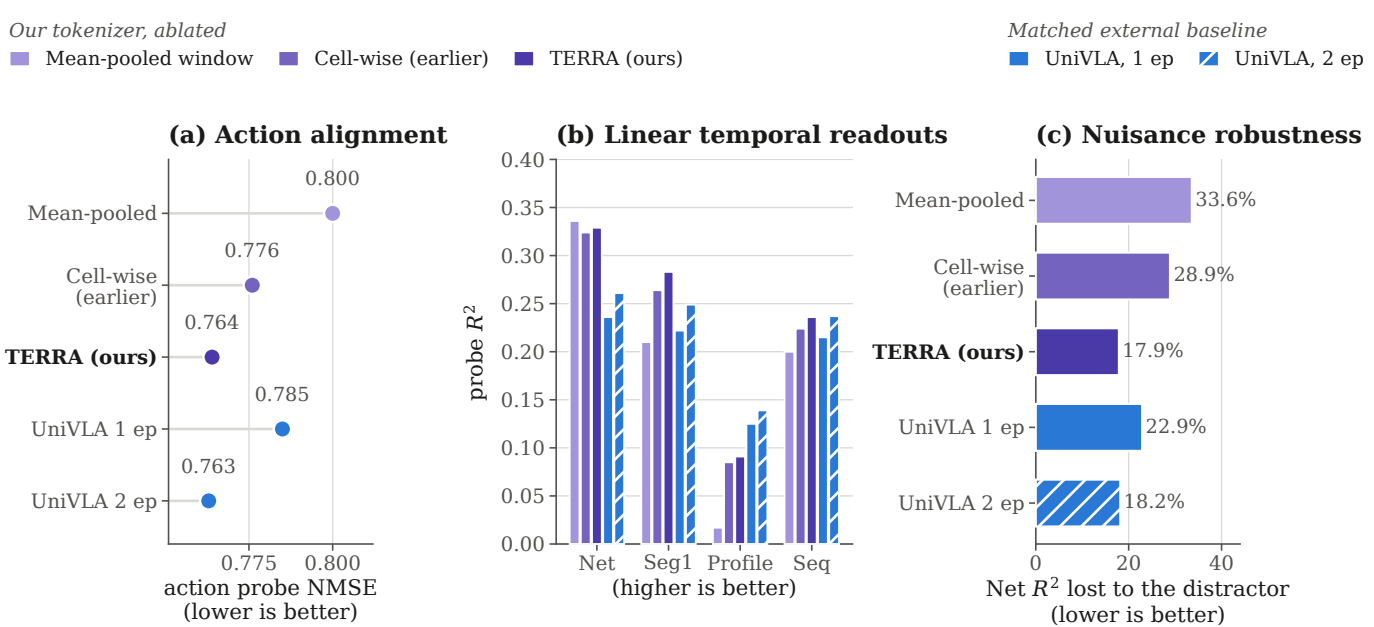
Our tokenizer, ablated
Mean-pooled window
Cell-wise (earlier)
TERRA (ours)
Matched external baseline
UniVLA, 1 ep
UniVLA, 2 ep
(a) Action alignment
0.800
Mean-pooled
0.776
Cell-wise (earlier)
0.764
TERRA (ours)
0.785
UniVLA 1 ep
0.763
UniVLA 2 ep
0.775
0.800
action probe NMSE (lower is better)
(b) Linear temporal readouts
probe R^2
0.40
0.35
0.30
0.25
0.20
0.15
0.10
0.05
0.00
Net
Seg1
Profile
Seq
(higher is better)
(c) Nuisance robustness
Mean-pooled
33.6%
Cell-wise (earlier)
28.9%
TERRA (ours)
17.9%
UniVLA 1 ep
22.9%
UniVLA 2 ep
18.2%
0
20
40
Net R^2 lost to the distractor (lower is better)